Getting started with Al - Computer vision 101

Michele Vitali

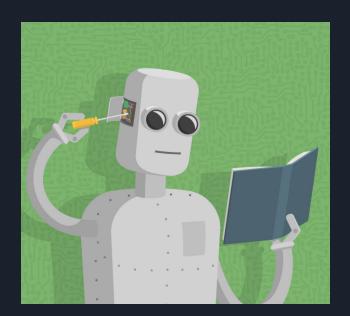
FabLab Parma, 27/10/2018

Agenda

Machine learning and Computer vision
Digit recognition - MNIST data set
Artificial neural networks
Fully connected neural networks
Deep neural networks
Convolutional neural networks

Machine learning

A particular field of computer science that gives computers the ability to learn tasks without being explicitly programmed

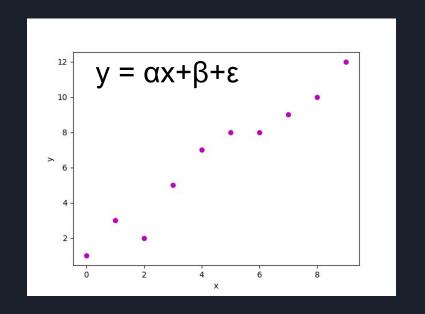


Machine learning

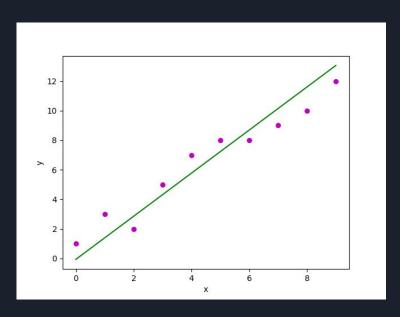
Learning from what?

X	у
0	1
1	3
2	2
3	5
4	7
5	8
6	8
7	9
8	10
9	12

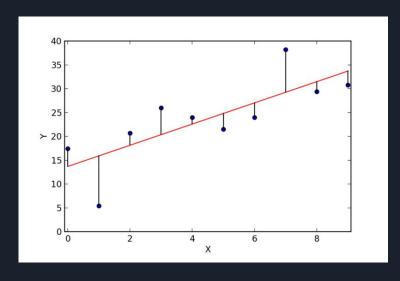
Data!



Model: y = Ax+B



The <u>best A and B</u> are those who <u>minimize the differences</u> between the model and the real data



Coefficients:

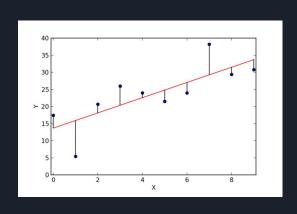
 $A \sim \alpha$

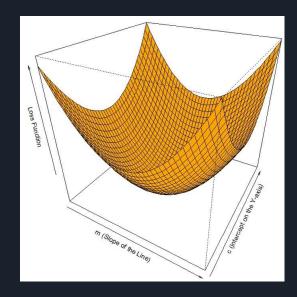
B ~ β

Loss function:

RMSE

The <u>best A and B</u> are those who <u>minimize the differences</u> between the model and the real data





Coefficients:

Α~α

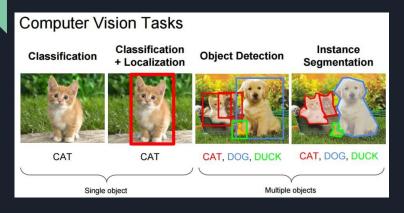
B ~ β

Loss function:

RMSE



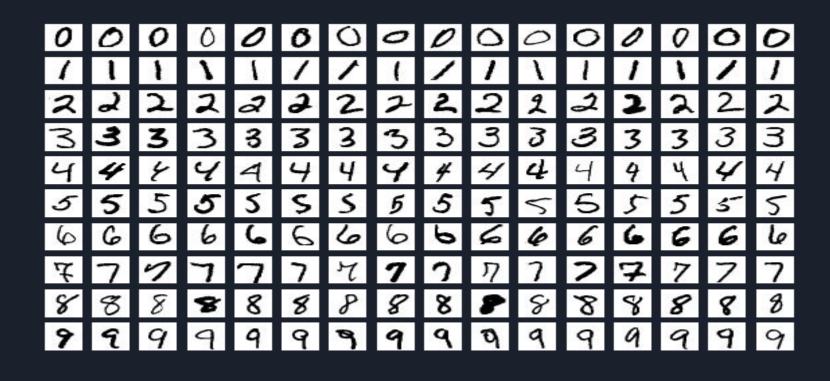
Computer vision



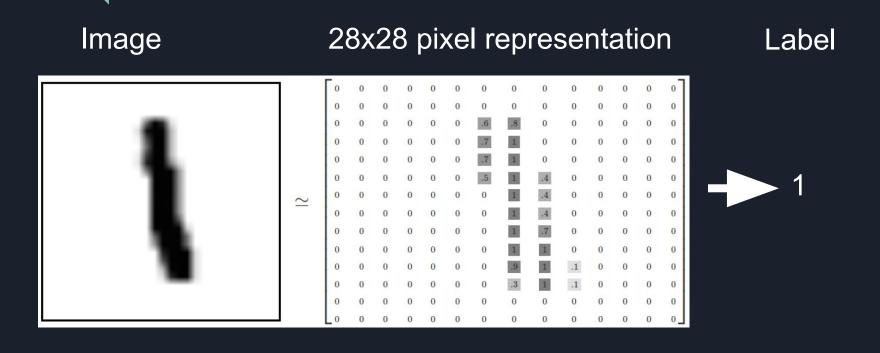




Digit recognition - MNIST dataset



MNIST dataset



MNIST dataset

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8		pixel774	pixel775	pixel776	pixel777	pixel778	pixel779	pixel780	pixel7
0	1	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	1
1	0	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
2	1	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
3	4	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
6	7	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
7	3	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
8	5	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
9	3	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
10	8	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
11	9	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
12	1	0	0	0	0	0	0	0	0	0	***	0	0	0	0	0	0	0	
13	3	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
14	3	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
15	1	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	
16	2	0	0	0	0	0	0	0	0	0		0	0	0	0	0	0	0	

Dataset preparation

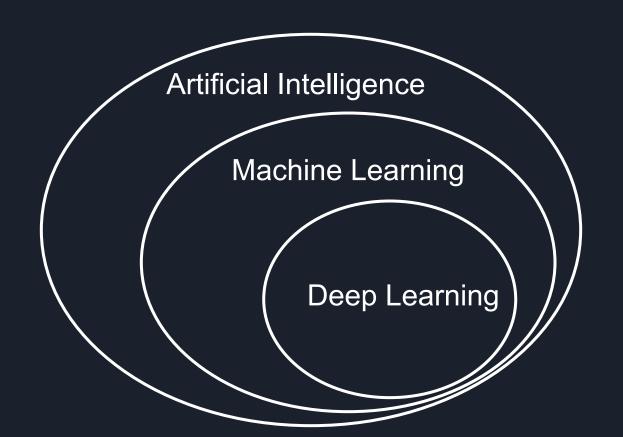
Main steps:

- ∞ Load dataset
- Split dataset in training set e validation set
- ∞ Prepare input (x) and output (y) data

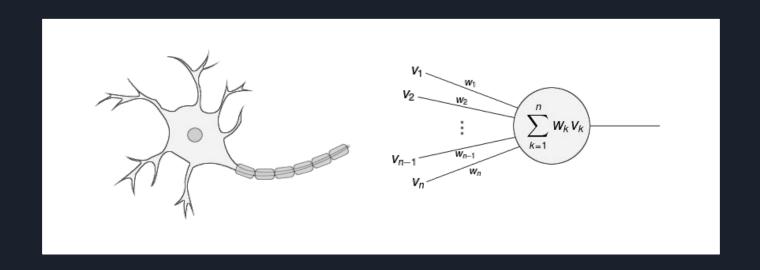
Dataset preparation



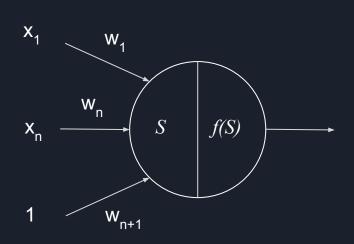
Deep learning



An artificial neural networks (ANN) are mathematical models inspired by biological neural networks

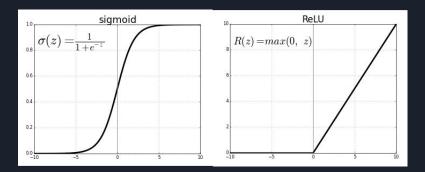


ANN are composed by artificial neurons

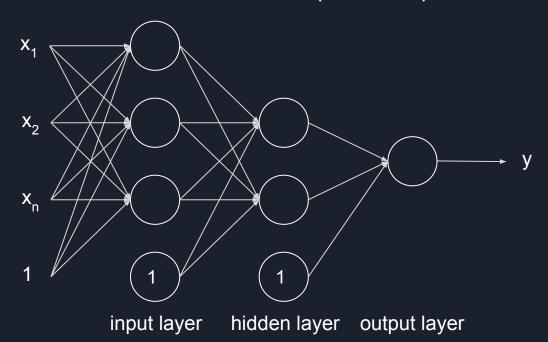


$$S = \sum x \cdot w$$

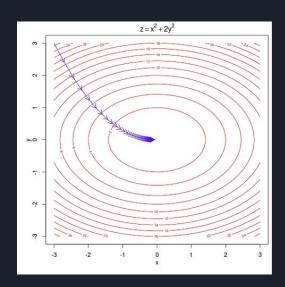
f(S)= Activation function: a non linear function applied to S

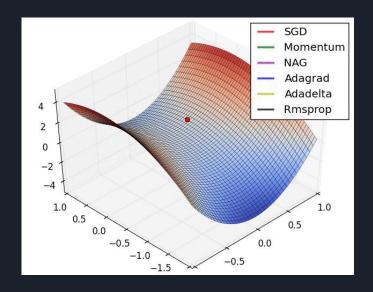


Fully connected neural network (shallow)



Gradient descent is used to obtain the set of weights that minimize the loss function





Don't worry, there are tools to implement ANNs! Some examples:



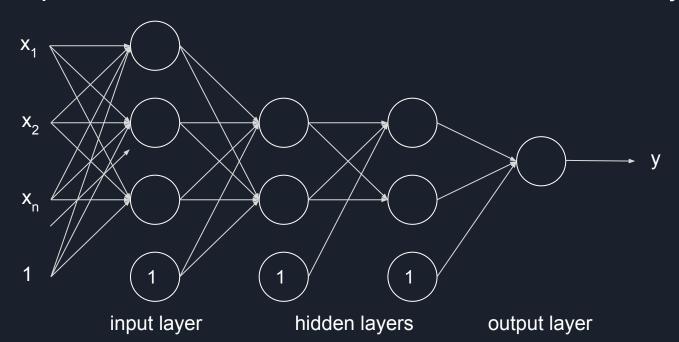


Fully connected neural networks (shallow)



Deep neural networks

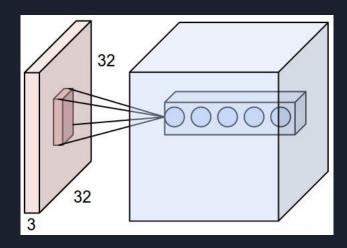
Deep neural networks have more than one hidden layer



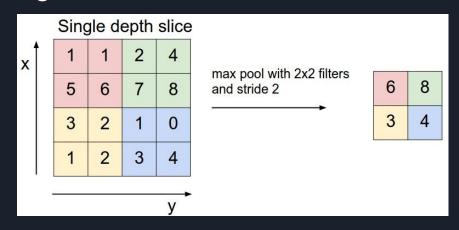
Deep neural networks



With data like images, it is impractical to connect neurons to all neurons in the previous layer. In CNN neurons are connected to only a local region of the input layer.



Pooling layers are inserted in-between successive Conv layers to progressively reduce the size of the representation and so the amount of parameters and computation to also control overfitting.



Stacking convolutional, pooling and fully connected layers creates the network architecture.

